


The PropPrompt Machine

An agent owns a territory, the wizard takes the brief, one of five AI platforms writes the analysis, and it leaves as a document. Everything below is what that sentence is actually made of — every page, entity, and backend function running in the Base44 app today, and what stops if a link breaks.



The one-line model, drawn: the territory is the license to run analyses at all, the engine does the work, and the finished analysis leaves by whichever door the agent picked.

The map

Eight systems inside one Base44 app. Every entity and function below is real and in the codebase today — grouped by the job it does, not by the folder it lives in.

The app

base44 · React + Vite

Base44 SDK

appId 69b849e8...

entities, auth, and functions.invoke

Layout.jsx

shell, nav, and the modal stack

25 pages

auto-registered from src/pages/

shadcn/ui + Radix

the whole component vocabulary

The territory market

what an agent actually buys

Territory / TerritorySubscription

one town or city, a fixed number of seats

PopulationPool

small towns combined into one bucket

Four ways to claim

single · pool · bundle · buyout

TerritoryBundle

many territories, one subscription

FullBuyoutSubscription

every seat in a city, zero competition

TerritoryClaimRequest

the claim, pending admin review

TerritoryWatchList

tell me when this one frees up

ReleasedTerritory

back on the map when a seat lapses

The analysis engine the product itself

6-step wizard

platform → assessment → client → property → format → confirm

Seven analysis types

listing pricing, buyer intel, CMA, portfolio, investment, rental, custom

Five stream functions

`claudeStream`, `chatgptStream`, `geminiStream`,
`perplexityStream`, `grokStream`

SSE, not invoke raw fetch

streaming needs the bare function URL

Public records

parcel data, with a missing-data question step

Three output formats

`narrative` · `structured` · `bullets`

checkAnalysisQuota

runs before the stream, not after

AIModelRouting

which model answers, set per platform

Delivery how the analysis leaves

generateDocuments

PDF and PPTX, branded per org

driveSync OAuth

straight into the agent's Google Drive

sendAnalysisEmail

to the client, from the app

crmPush / crmConnect

onto the deal record in their CRM

Org, identity & branding who sees whose logo

Organization / OrgMembership

brokerage → team → agent

SeatQuota

the monthly analysis cap, per org

Branding, three levels

brokerage, team, and agent each override

Members + invites

seats handed out inside the org

Money Stripe, and what it prices

stripeCheckout / createSetupIntent

the claim payment and the card on file

PricingConfig live

every price read at runtime, not hardcoded

TopupPack

10 / 25 / 50 / 100 analyses, one-time

RevenueShareLedger

what's owed to whoever brought the agent

PricingChangeLog

who changed a price, and when

CapExceededModal

the wall the agent hits at the cap

Compliance & training

the part that isn't optional

FairHousingReview

monthly, AI-generated, must be signed

signFairHousingReview

scroll to the end before the button works

logPrivacyEvent

the privacy notice, recorded

TrainingModule / Video / Progress

published modules, per-user completion

Admin

the console behind it

PlatformAdmin

prompts, model routing, platform config

ClaimsAdmin

approve or decline territory claims

PricingAdmin

the only place prices move

AdminDigest + DataQuality

what ran, and where the map is wrong

Three things moving through it

The map above is where things live. This is what actually travels between them, for the three journeys that matter.

An analysis, wizard to delivered

• New Analysis

agent picks the platform

→

• Six steps

assessment, client, property, format

→

• Quota check

checkAnalysisQuota, before any spend

→

• Streamed

SSE from one of five platforms

• Analysis record saved

readable at AnalysisDetail, private unless shared

→

• Out the door

PDF · PPTX · Drive · email · CRM

A territory, available to owned

• The map

Territories, filtered by status

→

• Claim type

single, pool, bundle, or buyout

→

• Tier + payment

Stripe, priced from PricingConfig

→

• Admin approves

ClaimsAdmin reviews the request

• Subscription active

seats_claimed rises; the town greys out on the map

→

• Analyses unlock

SeatQuota sets the monthly cap

A month, first analysis to the cap



If it breaks

The failure modes worth knowing before they happen.

A stream function goes down	Only that platform stops. The wizard still completes and the record still saves — the agent just can't generate on <code>claude</code> (or whichever) until it's back. Unset platforms fall back to <code>claudeStream</code> .
PricingConfig is wrong	Every price in the app is read live from it — claim tiers, top-up packs, upgrade prompts. One bad row misprices the whole funnel at once. PricingChangeLog is the only record of who moved it.
checkAnalysisQuota fails open	Analyses run past the cap and the org is never charged for them. The quota gate is the only thing standing between a subscription tier and unlimited AI spend.
Claims sit unapproved	Payment has already cleared through Stripe, but the territory stays available on the map and the agent has no seat. Nothing chases this on its own — it waits in ClaimsAdmin.
Fair Housing review unsigned	The overdue banner escalates, but the compliance exposure is the real cost. It is monthly, per user, and AI-generated — nobody notices a quiet month until an audit does.
The Modal host rotates	Dev-only, and now fixed: <code>server.allowedHosts</code> holds <code>.w.modal.host</code> as a suffix wildcard, so a regenerated sandbox URL no longer blocks the dev server.

Built from the PropPrompt codebase at `src/` — 25 pages, 38 entities, and 23 backend functions on Base44 app `69b849e8d5f86924955e7fae`. The code stays the source of truth; this page is the shape of it at a glance.